



On The Hardness of Learning to Stabilize Linear Systems

Xiong Zeng¹, Zexiang Liu¹, Zhe Du¹, Necmiye Ozay¹, Mario Sznaiar²

1: University of Michigan, Ann Arbor

2: Northeastern University, Boston



Introduction:

Learning an initial stabilizing controller is an important step in many learning-based control tasks. This poster studies the statistical hardness of learning to stabilize linear time-invariant systems. Our analysis shows a larger class of systems that are hard to learn to stabilize compared with earlier work [2]. This work is published in [1].

Problem Setup:

Given

- a discrete-time linear time-invariant (LTI) system S from a class \mathcal{C}_n :

$$x_{t+1} = Ax_t + Bu_t + w_t,$$

where $x_t \in \mathbb{R}^n$, $w_t \sim N(0, \sigma_w^2 I)$, $\sigma_w^2 > 0$,

- a time budget N and an input budget $\mathbb{E}[\|u_t\|^2] \leq \sigma_u^2$.

Consider a learning algorithm π that

- interacts with the system for N units of time, and
- outputs a linear static state-feedback controller \hat{K} at time N .

We want \hat{K} to stabilize the (unknown) system S .

Main Result: Consider two special LTI systems with

$$A_1 = A_2 = \begin{bmatrix} r & v & 0 & \dots & 0 \\ 0 & 0 & v & \dots & 0 \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & 0 & \dots & v \\ 0 & 0 & 0 & \dots & 0 \end{bmatrix}, B_1 = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ v \end{bmatrix}, B_2 = \begin{bmatrix} m \\ 0 \\ \vdots \\ v \end{bmatrix}.$$

Key insight: if the class \mathcal{C}_n contains a pair of systems (A_1, B_1) and (A_2, B_2) that are **hard to distinguish**, and **not co-stabilizable**, then **learning to stabilize is hard**.

Proposition 1: There exists a state-feedback controller K that stabilizes both (A_1, B_1) and (A_2, B_2) , only if

$$m \leq \mathcal{O}\left(\frac{1}{\exp(n)}\right).$$

Theorem 1: For any class \mathcal{C}_n of systems containing such a pair with $m = 2\left(\frac{2v}{r-1}\right)^n$ and any learning algorithm, the least amount of data for learning to stabilize arbitrary systems in \mathcal{C}_n **grows exponentially** with the state dimension n . Specifically,

$$\inf_{S \in \mathcal{C}_n} P_{S, \pi}^N(\rho(A + B\hat{K}) < 1) \geq 1 - \delta,$$

is satisfied only if

$$N \geq \frac{\delta_w^2}{2\delta_u^2} \left(\frac{r-1}{2v}\right)^{2n} \log \frac{1}{3\delta},$$

where $n \geq 2, r > 1$, and $0 < v < \frac{r-1}{2}$.

The key ideas in the proof:

- Use **Ackerman's formula** to show that (A_1, B_1) and (A_2, B_2) are not co-stabilizable (Proposition 1);
- Show that the KL divergence between (A_1, B_1) and (A_2, B_2) decays exponentially with n (indistinguishability).

i + ii + Birge's Inequality \rightarrow at least $\exp(n)$ samples to learn a stabilizing feedback gain for an unknown system being either (A_1, B_1) and (A_2, B_2) .

Comparison with Existing Result:

Tsiamis et al. [2]	Stable A	Degenerate noise
Our Work	Unstable A	Non-degenerate noise

Compared with [2], the system classes in our work are **not necessarily hard to identify**, but still **hard to stabilize**.



Experiments:

Certainty equivalence LQR

LMI-based sufficient condition for co-stabilizability

$$\min_{u_0, u_1, \dots, T \rightarrow \infty} \mathbb{E} \left[\frac{1}{T} \sum_{t=0}^T (x_t^T Q x_t + u_t^T R u_t) \right] \text{ find } K, P \text{ s.t. } \begin{cases} (A + B_1 K)^T P (A + B_1 K) < P \\ (A + B_2(m) K)^T P (A + B_2(m) K) < P \\ P > 0 \end{cases}$$

$$\text{s.t. } x_{t+1} = Ax_t + Bu_t + w_t$$




Current work:

Reduce sample complexity by using **dynamic time-varying** controllers.

References:

- [1] X. Zeng, Z. Liu, Z. Du, N. Ozay, and M. Sznaiar, On the Hardness of Learning to Stabilize Linear Systems, CDC, 2023
- [2] Tsiamis, Anastasios, et al. "Learning to control linear systems can be hard." Conference on Learning Theory. PMLR, 2022.



(Use your favorite algorithm π : model-free, model-based, etc.)

Question: Given a class \mathcal{C}_n of linear systems, how does the **sample complexity** for learning stabilizing static state-feedback controllers depend on the **system dimension n** ?