

# On the hardness of learning to stabilize linear systems

Xiong Zeng<sup>1</sup>



Zexiang Liu<sup>1</sup>



Zhe Du<sup>1</sup>



**Necmiye Ozay<sup>1</sup>**


Mario Sznaier<sup>2</sup>

<sup>1</sup>Electrical Engineering and Computer Science  
University of Michigan, Ann Arbor

<sup>2</sup>Electrical and Computer Engineering  
Northeastern University

# On the hardness of learning to stabilize linear systems with a static linear state-feedback controller

Xiong Zeng<sup>1</sup>



Zexiang Liu<sup>1</sup>



Zhe Du<sup>1</sup>

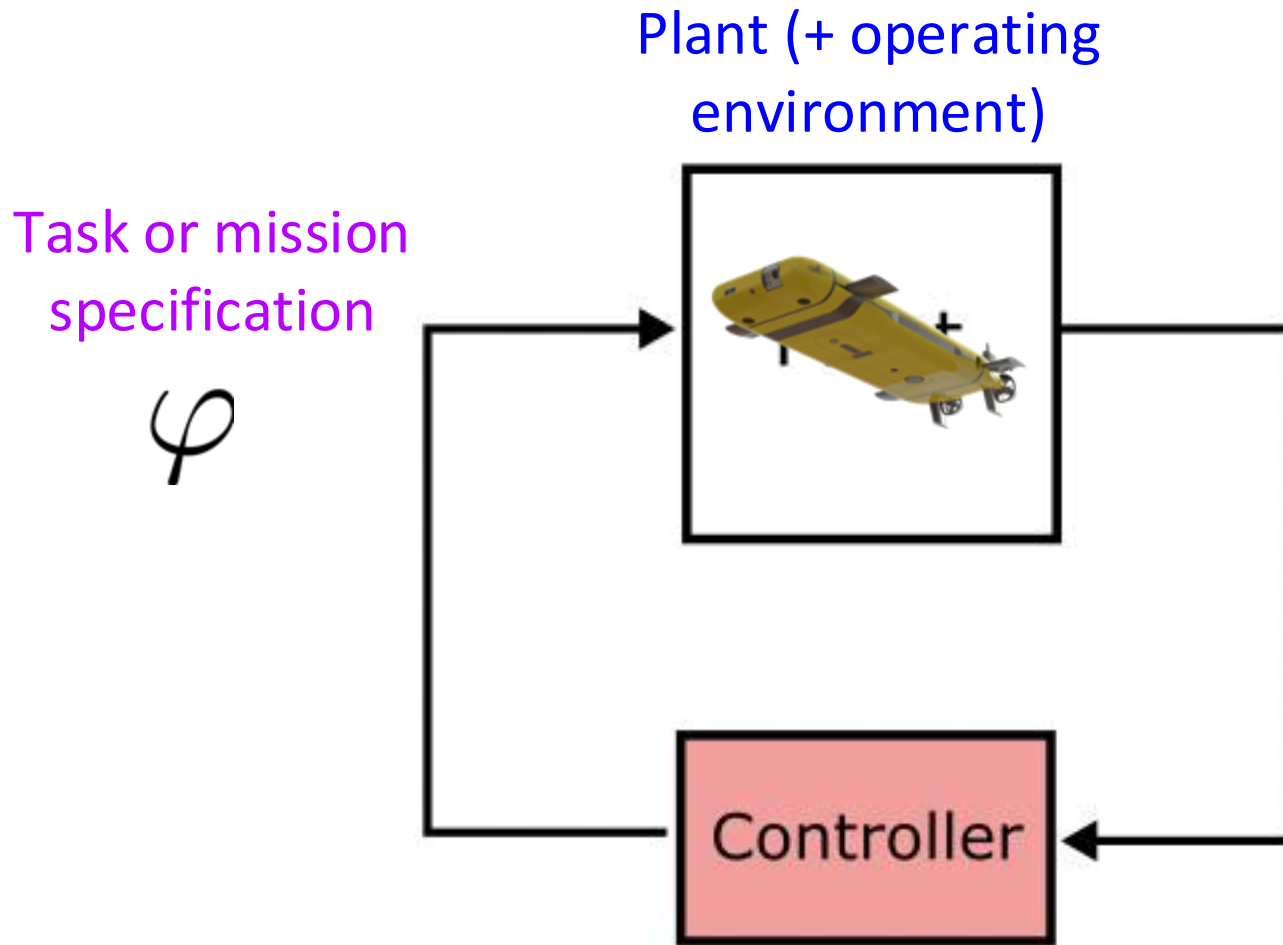


**Necmiye Ozay<sup>1</sup>**

Mario Sznaier<sup>2</sup>

<sup>1</sup>Electrical Engineering and Computer Science  
University of Michigan, Ann Arbor

<sup>2</sup>Electrical and Computer Engineering  
Northeastern University



Many components in a feedback loop can be learned:

- **Plant model:** model-based learning, system identification
- **Task specification:** inverse optimal control, inverse reinforcement learning
- **Controller:** model-free learning
- **Perception/estimation modules**

What are the fundamental limits in learning-based control?

# What is complexity?

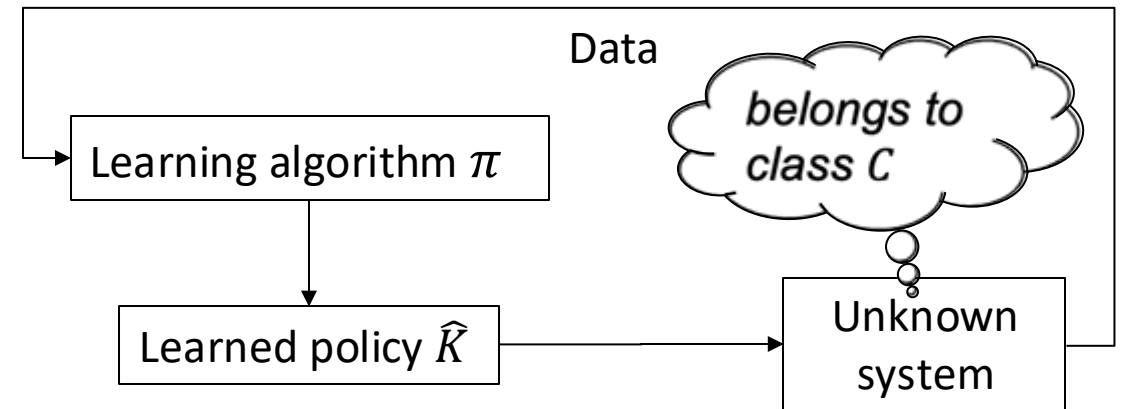
- Computational complexity: how the **computational cost** of an algorithm depends on the **input size**?

$$\text{Comp. Cost} = \mathcal{O}(\text{input size})$$

- Sample complexity: how a **performance metric** of a *learning* algorithm depends on the **sample size** (# of data points)?

$$\text{Performance} = \mathcal{O}(\text{sample size})$$

Examples of performance metric:  
approximation error,  
accumulative cost or regret,  
**probability of stability** (this work)



complexity of a problem  $\approx$  complexity of the best algorithm

# Setup

Given

- A discrete-time linear time-invariant (LTI) system  $S$  from a **class**  $\mathcal{C}_n$ :

$$x_{t+1} = Ax_t + Bu_t + w_t,$$

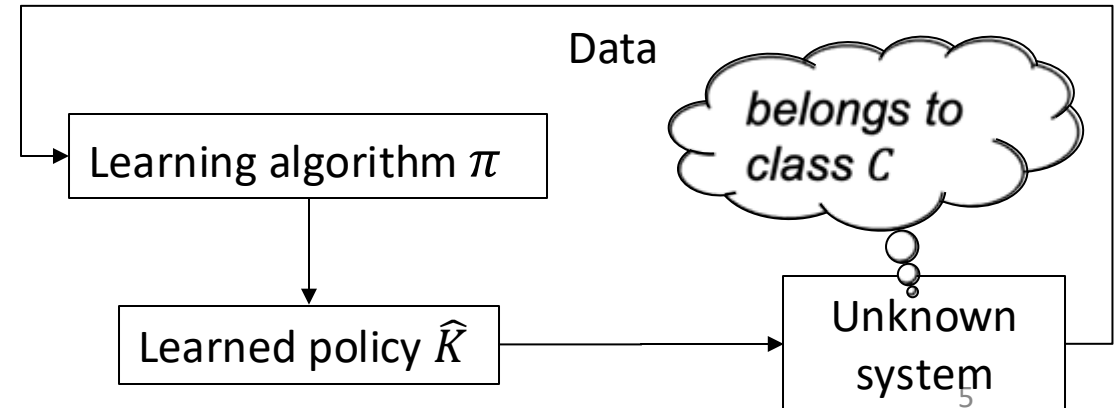
where  $x_t \in \mathbb{R}^n$  and  $w_t \sim \mathcal{N}(0, \sigma_w^2 \mathbf{I})$

- A time budget  $N$  and an input budget  $\mathbb{E}[\|u_t\|^2] \leq \sigma_u^2$

Consider a learning algorithm  $\pi$  that

- (i) interacts with the system for  $N$  units of time, and
- (ii) outputs a linear static state-feedback controller  $\hat{K}$  at time  $N$

We want  $\hat{K}$  to stabilize the (unknown) system  $S$ .



# Setup

Given

- A discrete-time linear time-invariant

where  $x_t \in \mathbb{R}^n$  and  $w_t \sim \mathcal{N}(0, \sigma_w^2 \mathbf{I})$

- A time budget  $N$  and an input budget

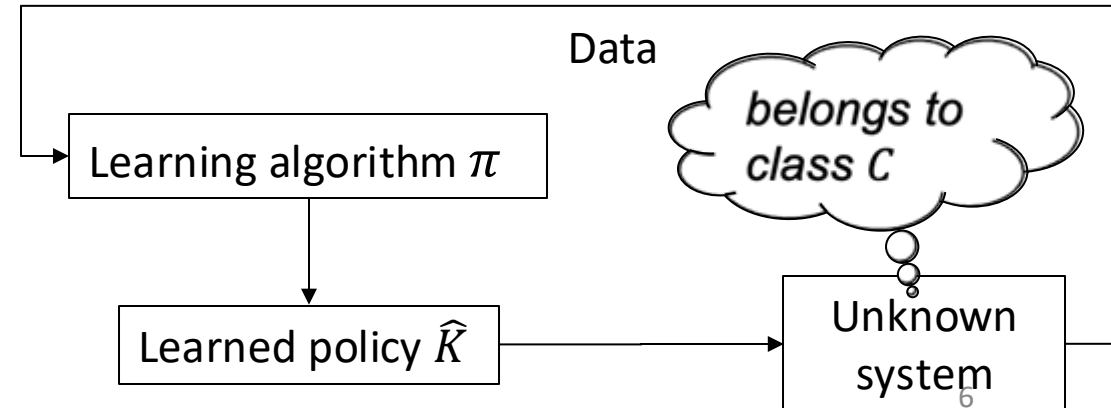
Use your favorite algorithm  $\pi$ :

- Excite the system for  $N$  steps and do (direct) data-driven control
- Excite the system for  $N$  steps and do system id and certainty equivalent control
- Use any other online/active learning method

Consider a learning algorithm  $\pi$  that

- (i) interacts with the system for  $N$  units of time, and
- (ii) outputs a linear static state-feedback controller  $\hat{K}$  at time  $N$

We want  $\hat{K}$  to stabilize the (unknown) system  $S$ .



# Problem statement

Denote the probability that  $\hat{K}$  stabilizes an unknown system  $S = (A, B) \in \mathcal{C}_n$  by

$$P_{S,\pi}^N \left( (A + B \hat{K}) \text{ is stable} \right)$$

A learning algorithm  $\pi$  is evaluated by the **worst-case** stabilization probability over  $\mathcal{C}_n$ :

$$\inf_{S \in \mathcal{C}_n} P_{S,\pi}^N \left( (A + B \hat{K}) \text{ is stable} \right)$$

**Sample complexity** of  $\pi$ : Given a  $\delta \in [0,1]$ , the **minimum time/data budget  $N$**  such that

$$\inf_{S \in \mathcal{C}_n} P_{S,\pi}^N \left( (A + B \hat{K}) \text{ is stable} \right) \geq 1 - \delta.$$

**Question:** Given a class  $\mathcal{C}$  of linear systems, how does the **sample complexity** for learning stabilizing static state-feedback controllers depend on the **system dimension  $n$** ?

# Related work

**Question:** Given a class  $\mathcal{C}$  of linear systems, how does the **sample complexity** for learning stabilizing static state-feedback controllers depend on the **system dimension  $n$** ?

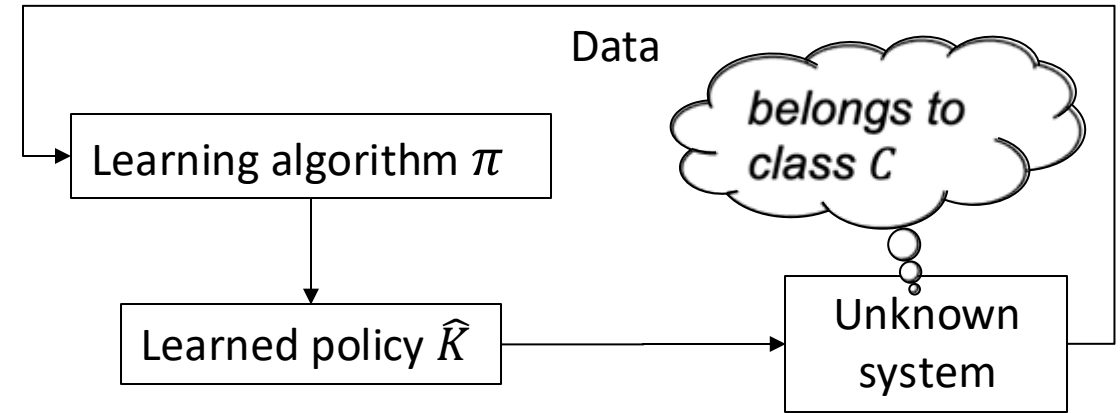
**Existing results (very incomplete list):**

Dean et al. 2017, Mania et al. 2019	$N \geq \mathcal{C}^{sys} poly(n)$ $\Rightarrow Stable$	Non-degenerate noise ( $w_t \sim N(0, \Sigma), \Sigma > 0$ )
Tsiamis et al. 2022	$Stable \Rightarrow$ $N \geq O(exp(n))$	Degenerate noise ( $\Sigma \geq 0$ and $\Sigma \neq 0$ ): <b>Hard to identify</b>
<b>Our Work</b>	$Stable \Rightarrow$ $N \geq O(exp(n))$	Non-degenerate noise ( $\Sigma > 0$ ): <b>Not necessarily hard to identify but there is a different obstruction</b>

**Sufficient** conditions  
for sample complexity

**Necessary** conditions  
for sample complexity

aka, **lower bounds**





# Roadmap to establish hardness

Find a class  $\mathcal{C}_n$  of systems such that for any algorithm  $\pi$  and for all  $\delta \in [0,0.5]$ , we have

$$\inf_{S \in \mathcal{C}_n} P_{S,\pi}^N((A + B \hat{K}) \text{ is stable}) \geq 1 - \delta.$$

only if

$$N \geq O(\exp(n)).$$

**Key insight:** if the class  $\mathcal{C}_n$  contains a pair of systems  $(A_1, B_1)$  and  $(A_2, B_2)$  that are

- **hard to distinguish**, and
- **not co-stabilizable**

**then learning to stabilize is hard.**

# Roadmap to establish hardness

Find a class  $\mathcal{C}_n$  of systems such that for any algorithm  $\pi$  and for all  $\delta \in [0,0.5]$ , we have

$$\inf_{S \in \mathcal{C}_n} P_{S,\pi}^N((A + B \hat{K}) \text{ is stable}) \geq 1 - \delta.$$

only if

$$N \geq O(\exp(n)).$$

**Key insight:** if the class  $\mathcal{C}_n$  contains a pair of systems  $(A_1, B_1)$  and  $(A_2, B_2)$  that are

- **hard to distinguish**, and  
when driven by the same policy  $\pi$ , the resulting state trajectories are similar
- **not co-stabilizable**  
there does not exist a single controller that simultaneously stabilizes the two systems

**then learning to stabilize is hard.**  $\nexists K$  s.t. both  $A_1 + B_1 K$  and  $A_2 + B_2 K$  are stable

## A pair of systems that are hard to learn to stabilize

$$\mathbf{A} = \begin{bmatrix} r & v & 0 & \cdots & 0 \\ 0 & 0 & v & \cdots & 0 \\ & & \ddots & \ddots & \\ 0 & 0 & 0 & \cdots & v \\ 0 & 0 & 0 & \cdots & 0 \end{bmatrix} \in \mathbb{R}^{n \times n}, \mathbf{B} = \begin{bmatrix} b^{(1)} \\ 0 \\ \vdots \\ 0 \\ v \end{bmatrix} \in \mathbb{R}^n$$

$x_{t+1} = Ax_t + Bu_t + w_t, \quad r > 1, 0 < v < \frac{r-1}{2}, \text{ and } b^{(1)} \geq 0.$

First order system with  
actuation delay of length  $n-1$

Take  $B_1 = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ v \end{bmatrix}, B_2 = \begin{bmatrix} m \\ 0 \\ \vdots \\ v \end{bmatrix}$

# A pair of systems that are hard to learn to stabilize

$$A = \begin{bmatrix} r & v & 0 & \cdots & 0 \\ 0 & 0 & v & \cdots & 0 \\ & & \ddots & \ddots & \\ 0 & 0 & 0 & \cdots & v \\ 0 & 0 & 0 & \cdots & 0 \end{bmatrix} \in \mathbb{R}^{n \times n}, \quad B = \begin{bmatrix} b^{(1)} \\ 0 \\ \vdots \\ 0 \\ v \end{bmatrix} \in \mathbb{R}^n$$

$x_{t+1} = Ax_t + Bu_t + w_t, \quad r > 1, 0 < v < \frac{r-1}{2}, \text{ and } b^{(1)} \geq 0.$

First order system with  
actuation delay of length  $n-1$

Take  $B_1 = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ v \end{bmatrix}, B_2 = \begin{bmatrix} m \\ 0 \\ \vdots \\ v \end{bmatrix}$

**Proposition 1:** There exists a state-feedback controller  $K$  that stabilizes both  $(A, B_1)$  and  $(A, B_2)$  only if

$$m \leq \mathcal{O}\left(\frac{1}{\exp(n)}\right)$$

## A pair of systems that are hard to learn to stabilize

**Proposition 1:** There exists a state-feedback controller  $K$  that stabilizes both  $(A, B_1)$  and  $(A, B_2)$  only if

$$m \leq \mathcal{O}\left(\frac{1}{\exp(n)}\right)$$

*Proof sketch:* Use Ackerman's formula to analytically construct  $K$  to place the poles of system 1 arbitrarily inside the unit circle and check the stability of system 2 when controlled by  $K$  using Jury stability test.

# A pair of systems that are hard to learn to stabilize


**Proposition 1:** There exists a state-feedback controller  $K$  that stabilizes both  $(A, B_1)$  and  $(A, B_2)$  only if

$$m \leq \mathcal{O}\left(\frac{1}{\exp(n)}\right)$$

*Proof sketch:* Use Ackerman's formula to analytically construct  $K$  to place the poles of system 1 arbitrarily inside the unit circle and check the stability of system 2 when controlled by  $K$  using Jury stability test.

We can also do a sanity check numerically using LMI-based sufficient condition for co-stabilizability


$$\begin{aligned} & \text{find } K, P \\ & \text{s.t. } (A + B_1 K)^T P (A + B_1 K) < P \\ & \quad (A + B_2(m) K)^T P (A + B_2(m) K) < P \\ & \quad P > 0 \end{aligned}$$



# Main result

$$\mathbf{A} = \begin{bmatrix} r & v & 0 & \cdots & 0 \\ 0 & 0 & v & \cdots & 0 \\ & & \ddots & \ddots & \\ 0 & 0 & 0 & \cdots & v \\ 0 & 0 & 0 & \cdots & 0 \end{bmatrix} \in \mathbb{R}^{n \times n}, \mathbf{B} = \begin{bmatrix} b^{(1)} \\ 0 \\ \vdots \\ 0 \\ v \end{bmatrix} \in \mathbb{R}^n$$

$$x_{t+1} = Ax_t + Bu_t + w_t, \quad r > 1, 0 < v < \frac{r-1}{2}, \text{ and } b^{(1)} \geq 0.$$



$$\text{Take } B_1 = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ v \end{bmatrix}, B_2 = \begin{bmatrix} m \\ 0 \\ \vdots \\ v \end{bmatrix}$$

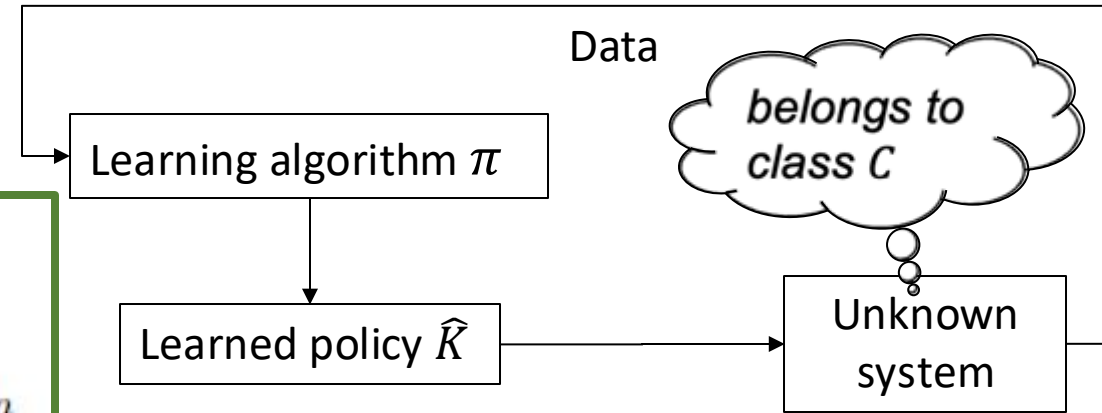
**Theorem:** For **any** class  $\mathcal{C}$  of systems containing such a pair with  $m = 2 \left( \frac{2v}{r-1} \right)^n$  and **any** learning algorithm, the **least amount of data** for learning to stabilize arbitrary systems in  $\mathcal{C}$  **grows exponentially** with the state dimension!

$$\inf_{S \in \mathcal{C}_n} P_{S, \pi}^N((A + B \hat{K}) \text{ is stable}) \geq 1 - \delta \quad \text{only if} \quad N \geq \underbrace{\frac{\sigma_w^2}{2\sigma_u^2} \left( \frac{r-1}{2v} \right)^{2n}}_{\text{SNR}} \log\left(\frac{1}{3\delta}\right) \rightarrow \text{exp in dimension } n$$

# Main result

$$\mathbf{A} = \begin{bmatrix} r & v & 0 & \cdots & 0 \\ 0 & 0 & v & \cdots & 0 \\ & & \ddots & \ddots & \\ 0 & 0 & 0 & \cdots & v \\ 0 & 0 & 0 & \cdots & 0 \end{bmatrix} \in \mathbb{R}^{n \times n}, \mathbf{B} = \begin{bmatrix} b^{(1)} \\ 0 \\ \vdots \\ 0 \\ v \end{bmatrix} \in \mathbb{R}^n$$

$$x_{t+1} = Ax_t + Bu_t + w_t, \quad r > 1, 0 < v < \frac{r-1}{2}, \text{ and } b^{(1)} \geq 0.$$



$$\text{Take } B_1 = \begin{bmatrix} 0 \\ 0 \\ \vdots \\ v \end{bmatrix}, B_2 = \begin{bmatrix} m \\ 0 \\ \vdots \\ v \end{bmatrix}$$

**Theorem:** For **any** class  $\mathcal{C}$  of systems containing such a pair with  $m = 2 \left( \frac{2v}{r-1} \right)^n$  and **any** learning algorithm, the **least amount of data** for learning to stabilize arbitrary systems in  $\mathcal{C}$  **grows exponentially** with the state dimension!

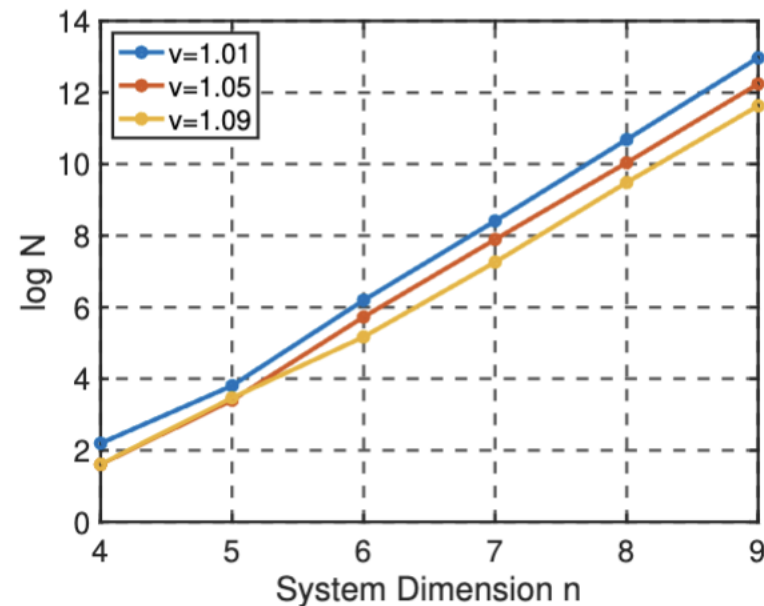
## Proof Sketch

- Having  $b^{(1)}$  entries of the pairs exponentially closer to each other as a function of  $n$  upper bounds the KL divergence of length- $N$  trajectories generated by the two systems (**indistinguishability**)
- **Co-stabilizability** requires exponential closeness in parameters space due to Proposition 1
- Use Birge's inequality to establish a KL divergence lower bound



# Numerical illustration

Number of samples required for 90% probability of finding a stabilizing controller using certainty equivalent LQR\*:



y-axis: observed **least sample size**  $N$  for CE LQR being stabilizing with probability 0.9

## Certainty Equivalent LQR

- Excite the system with iid Gaussian input for  $N$  steps
- Learn system matrices
- Do LQR with learned model

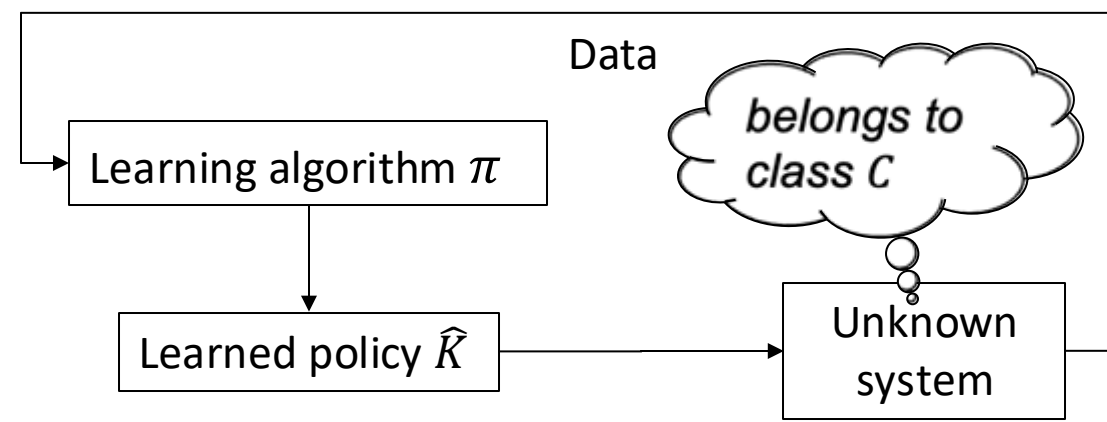
$$\min_{u_0, u_1, \dots} \lim_{T \rightarrow \infty} E \left[ \frac{1}{T} \sum_{t=0}^T (x_t^T Q x_t + u_t^T R u_t) \right]$$
$$s. t. x_{t+1} = A x_t + B u_t + w_t$$

\* By our result, same trend holds for any algorithm

# Conclusions

## Key contributions:

- Learning to stabilize a linear system is hard if we have systems that are hard to **distinguish** **AND** hard to **co-stabilize**
- Compare with earlier work by Tsiamis et al., our analysis
  - reveals a new obstruction for the hardness of learning to stabilize;
  - covers a larger class of systems (including open-loop unstable ones)



# Conclusions

## Key contributions: *with a static linear state-feedback controller*

- Learning to stabilize a linear system is hard if we have systems that are hard to **distinguish** **AND** hard to **co-stabilize**
- Compare with earlier work by Tsiamis et al., our analysis
  - reveals a new obstruction for the hardness of learning to stabilize;
  - covers a larger class of systems (including open-loop unstable ones)

## Current work:

- Observation: Static linear state feedback controllers are not a rich enough class for co-stabilizing linear systems (memory helps!)
- Can we reduce sample complexity by using dynamic time-varying controllers?

